

ECEC 414

Simulation Assignment – Introduction to gem5

Due January 22nd 2014

1 Overview

This quarter you will be studying concepts related to parallel architecture design and evaluation. You will later be given a practical simulation based assignment that tests those new concepts. Consequently, you should be proficient with using a processor simulator. This assignment is intended to be a gentle introduction to processor simulation.

2 Getting Started

In this class, we will use the GEM5 simulator. To work with the GEM5 simulator, you will need a set of scripts and configuration files. These can be copied to your home directory and decompressed with the following simple steps:

1. Login to `xunil.coe.drexel.edu` via SSH with your Drexel user id and password. Note: You should use either the linux tool “screen” or computer with a reliable Internet connection; both build and simulation tasks can take over an hour.
2. A stable version of the GEM5 simulator can be copied from my home directory. Use the following command:

```
> cp /home/DREXEL/sn446/gem5-stable.tar.bz2 .
```
3. Unpack the archive with the command:

```
> tar -xjvf gem5-stable.tar.bz2
```

 (execute `man tar` to get the manual-page for an explanation of the `tar` command)

You will now have a `gem5-stable` directory located in your home directory that contains the GEM5 simulator.
4. Navigate to the `gem5-stable` directory.

```
> cd gem5-stable
```

3 Building an Architecture

GEM5 is a robust architecture simulator that allows you to model many different ISAs, including ALPHA, ARM, POWER, MIPS, and x86. For many of these ISAs, you can configure either in-order and out-of-order (OoO or O3) processor models. The list of available ISAs is in the `gem5-stable/build` directory.

For this assignment, we will be building an out-of-order ALPHA processor. From your `gem5-stable` directory, execute the following command:

```
> scons build/ALPHA/gem5.opt
```

`scons` is similar to `make`, and compiles the code for the simulator, building the `gem5.opt` executable in the `~/gem5-stable/build/ALPHA` directory

4 Basic Usage

Now we will get started with the GEM5 simulator through a very basic usage example. Here, we will execute a set of benchmarks on GEM5 called Splash2. Splash2 contains several standard algorithms: Cholesky, FFT, LUContig, LUNoncontig, Radix, Barnes, FFM, OceanContig, OceanNoncontig, Raytrace, WaterNSquared, and WaterSpatial.

Here are the steps to run a splash-2 benchmark on GEM5:

1. Navigate to your home directory using the following command:

```
> cd ~
```
2. Download the splash-2 benchmarks using the following command:

```
> wget http://www.gem5.org/dist/m5_benchmarks/v1-splash-alpha.tgz
```
3. Extract the splash-2 benchmarks:

```
> tar -xzf v1-splash-alpha.tgz
```

You now have a directory `~/splash2` containing the benchmarks
4. Navigate to your build directory

```
> cd ~/gem5-stable/build/ALPHA
```
5. Copy the `runSplash.py` file from my folder to your build directory

```
> cp /home/DREXEL/sn446/runSplash.py .
```
6. Invoke the GEM5 simulator with the following command line:

```
> ./gem5.opt runSplash.py -d -n 4 --l1size 32kB --cacheblocksize 64 -b FFT  
--rootdir=/home/DREXEL/< drexel_id >/splash2/codes
```

`--l1size` configures the L1 cache size `--cacheblocksize` configures the cache line size, the amount of data transferred from memory on a cache miss. `-d` tells GEM5 to use the detailed out-of-order model.

`-n 4` tells GEM5 to create a system with 4 CPUs and invokes the benchmark with 4 threads

`-b FFT` tells GEM5 to run the FFT benchmark

`--rootdir` tells GEM5 where to find the benchmarks, replace `< drexel_id >` with your username

The simulator will run and generate output in the directory `./m5out`. Here you will find `config.ini` containing the simulation parameters and `stats.txt` containing the results of the simulation.

5 Cache Experiment

The goal of this experiment is to evaluate how cache configurations affect benchmark performance. You will use GEM5 to investigate cache miss rates, bus traffic, and execution time under different system configurations. You will be provided with a base cache/processor configuration and you will perform a set of experiments designed to test the effect of scalability and cache capacity. Pick one of the following experiments:

- keep L1 cache capacity fixed at 32KB, block size fixed at 64B and vary number of processors (2, 4, 8)
- keep number of processors fixed at 4, block size fixed at 64B and vary L1 cache capacity (16KB, 32KB, 64KB)

- keep L1 cache capacity fixed at 32KB, number of processors fixed at 4 and vary L1 cache block size (32B, 64B, 128B)

You should use the default cache coherence protocol (MSI), L2, and memory configurations. Pick any two benchmarks of your choosing out of the Splash2 suite: **Cholesky**, **FFT**, **LUContig**, **LUNoncontig**, **Radix**, **Barnes**, **FFM**, **OceanContig**, **OceanNoncontig**, **Raytrace**, **WaterNSquared**, or **WaterSpatial**. You should submit a brief assignment report which explains your configuration and reports:

- Data L1 cache miss rate (in each experiment, use the avg over all processors)
- L2 bus traffic (transactions per processor)
- Total performance (execution cycles), and IPC

You should explain how varying the parameters impacts your results. Then offer your guess as to why you see the results you do. What do the results tell you about your workloads? Plot IPC, miss rate, bus traffic, and execution cycles for each benchmark across your configurations.